


Deep Learning for NLP

Large Language Models (LLMs)

Post-Training: Alignment and Reasoning

Learning goals

- understand supervised fine-tuning (SFT) and instruction tuning
- understand preference learning with reward models, PPO, and DPO
- understand behavioral trade-offs such as reward hacking and over-refusal
- learn about chain-of-thought reasoning



PRETRAINING DOES NOT PRODUCE AN ASSISTANT

A pretrained language model learns to predict the next token:

$$p_{\theta}(x_t \mid x_{<t})$$

- It learns language, knowledge, and many useful capabilities.
- But it is not explicitly trained to **follow a user's instruction**.
- Given a prompt, it may continue the text rather than answer the request.

► OpenAI, 2022

Prompt

Explain the moon landing to a 6 year old in a few sentences.

Completion

GPT-3

Explain the theory of gravity to a 6 year old.

Explain the theory of relativity to a 6 year old in a few sentences.

Explain the big bang theory to a 6 year old.

Explain evolution to a 6 year old.

Predicting text is not the same objective as being a useful assistant.

FROM PRETRAINING TO POST-TRAINING

After pretraining, additional training is used to shape model behavior.

Pretrained LM → SFT → Preference / RL training → Assistant

- **SFT:** learn from examples of desired responses.
- **Preference optimization:** learn which responses are preferred.
- Reasoning and safety can also be shaped during post-training.

Post-training changes how the capabilities learned during pretraining are used.

Supervised Fine-Tuning

WHAT IS SUPERVISED FINE-TUNING?

Supervised fine-tuning (SFT) trains the pretrained model on **instruction–response pairs**.

$$(x, y) = (\text{instruction / context}, \text{desired response})$$

For example:

User: Explain why the sky is blue.

Assistant: Sunlight is scattered by molecules in the atmosphere ...

- The desired response is provided by humans or another model.
- The model parameters are updated.
- This is therefore **training**, not prompting.

SFT USES THE LANGUAGE-MODEL OBJECTIVE

The model is still trained by next-token prediction.

For instruction x and desired response $y = (y_1, \dots, y_T)$:

$$\mathcal{L}_{\text{SFT}} = - \sum_{t=1}^T \log p_{\theta}(y_t \mid x, y_{<t}).$$

- The architecture does not need to change.
- What changes is the **training data**.
- Instead of arbitrary web text, the model sees examples of the behavior we want.


Same prediction objective, different data distribution.

WHICH TOKENS CONTRIBUTE TO THE SFT LOSS?

A training sequence contains both a user prompt and an assistant response.

```
<user> Explain why the sky is blue. </user>  
<assistant> Sunlight is scattered ...</assistant>
```

Typically,



- The prompt is given to the model as context.
- The loss is typically computed on the assistant response.
- The model learns which response should follow a given instruction.

WHERE DOES SFT DATA COME FROM?

In InstructGPT, SFT examples were built from prompts and human-written demonstrations.

Prompt → Human writes a good response → SFT example


- Prompts were largely drawn from the distribution of real user requests.
- Initially, annotators also wrote prompts to bootstrap the dataset.
- The main challenge is obtaining **high-quality desired responses**.

► Ouyang et al., 2022: Training language models to follow instructions with human feedback


FLAN: AN EARLY EXAMPLE OF INSTRUCTION TUNING

FLAN (Fine-tuned LAnguage Net) fine-tunes a pretrained language model on many tasks written as natural-language instructions.

(A) Pretrain–finetune (BERT, T5)



(B) Prompting (GPT-3)



(C) Instruction tuning (FLAN)





Figure 2: Comparing instruction tuning with pretrain–finetune and prompting.


- Tasks are converted into prompted natural-language formats.
- One model is trained across many different tasks.

HOW DOES FLAN FINE-TUNING WORK?

Finetune on many tasks (“instruction-tuning”)




Inference on unseen task type



- Each task is represented through natural-language instructions.
- The model is fine-tuned jointly on many such tasks.
- At test time, it can be given an instruction for a task that was not used directly for training.

INSTRUCTION TUNING GENERALIZES TO NEW TASKS

A central result of FLAN is improved performance on **held-out tasks**.



- The evaluation tasks were not included in the fine-tuning mixture.
- The instruction tells the model what task to perform.
- Instruction tuning therefore improves zero-shot task generalization.

► Wei et al., 2021

WHAT CAN SFT NOT EASILY SPECIFY?

Suppose several responses to the same prompt are all plausible.

Response A

plausible answer

Response B

another plausible answer

- Writing the single “ideal” response can be difficult or ill-defined.
- Many desirable properties are hard to express through one target sequence.
- It may be much easier for a human to answer:

Which response is better?

This motivates **preference learning**.

Preference Learning and RLHF

MOTIVATION: BACKFLIPPING

- What a backflip is: [▶ backflip](#)
- How to teach a backflip?

ORIGIN OF RLHF: LEARN HOW TO BACKFLIP

How to label data for training a backflipper?

- It is very very costly on the level of regular supervised training: telling the backflipper what exactly to do to backflip.
- Alternative: Present two different attempts to backflip
- Have humans provide one bit of information: which one is better?
- This animation shows what we want to learn: ▶ trained backflipper
- Human chooses one of two clips = one bit: ▶ example of human feedback
- Trained backflipper: ▶ trained backflipper

PREFERENCE DATA

Instead of asking humans to write the ideal response, we can ask them to **compare model-generated responses**.

For a prompt x , suppose a human prefers y^+ over y^- :

$$(x, y^+, y^-).$$

- Several responses can be generated for the same prompt.
- Human labelers rank the responses by quality.

Ranking outputs				
To be ranked				
1 A team of researchers from Yale University and University of California, Davis studied the vocalization patterns of several different types of parrots. They found that parrots like to mimic human speech, and can produce a wide range of sounds, such as whistles, squawks, and other types of vocalizations...	2 Parrots have been found to have the ability to understand numbers. Researchers have found that parrots can understand numbers up to six. In a series of experiments, the parrots were able to identify the amount of food items under a number of cups...			
Rank 1 (best)	Rank 2	Rank 3	Rank 4	Rank 5 (worst)
1 A research group in the United States has found that parrots can imitate human speech with ease, and some of them can even do so in the same way as humans. This group studied the sounds that parrots make in their natural habitats and found that they use their tongues and beaks in ways that are strikingly...		2 Scientists have found that green-winged parrots can tell the difference between two noises that are the same except for the order in which they are heard. This is important because green-winged parrots are known to imitate sounds. This research shows that they are able to understand the difference between sounds.		
		3 Current research suggests that parrots see and hear things in a different way than humans do. While humans see a rainbow of colors, parrots only see shades of red and green. Parrots can also see ultraviolet light, which is invisible to humans. Many birds have this ability to see ultraviolet light, an ability...		

► Ouyang et al., 2022: InstructGPT

WHAT DOES “BETTER” MEAN?

In InstructGPT, labelers were asked to prefer responses that are

helpful, truthful, and harmless.

But these objectives can conflict.


- A very helpful answer may contain incorrect information.
- A harmless answer may refuse a request that could have been answered safely.
- Different people may disagree about which response is preferable.

Preference labels encode judgments about desired model behavior.

► Ouyang et al., 2022: InstructGPT

LEARNING FROM PREFERENCES

The original InstructGPT pipeline has three training stages:



- **SFT.** Training data: human-written responses.
- **Reward model.** Training data: human rankings of model-generated responses.
- **Proximal Policy Optimization (PPO).** Optimization signal: trained reward model.

WHAT IS A REWARD MODEL?

A reward model assigns a scalar score to a prompt–response pair:

$$r_{\theta}(x, y) \in \mathbb{R}.$$

prompt x + response $y \longrightarrow$ reward model $\longrightarrow$ score $r_{\theta}(x, y)$

- Higher reward should correspond to responses humans prefer.
- θ denotes the parameters of the reward model.
- In InstructGPT, the reward model is trained from human-ranked model outputs.
- Same architecture as the SFT model, but smaller: InstructGPT used a 6B reward model when the policy had 175B parameters.

$$\text{loss}(\theta) = -\frac{1}{\binom{K}{2}} E_{(x, y_w, y_l) \sim D} [\log(\sigma(r_{\theta}(x, y_w) - r_{\theta}(x, y_l)))]$$

LEARNING THE REWARD MODEL FROM PREFERENCES

Suppose the labeler prefers y_w over y_l .

We want

$$r_{\theta}(x, y_w) > r_{\theta}(x, y_l).$$

In InstructGPT, the reward-model loss is

$$\text{loss}(\theta) = -\frac{1}{\binom{K}{2}} \mathbb{E}_{(x, y_w, y_l) \sim D} [\log \sigma(r_{\theta}(x, y_w) - r_{\theta}(x, y_l))].$$

- x : prompt.
- y_w : preferred (winning) response.
- y_l : lower-ranked response.
- D : fixed dataset of human preference comparisons.
- A ranking of K responses yields $\binom{K}{2}$ pairwise comparisons.

WHAT HAS THE REWARD MODEL LEARNED?

The reward model does **not** directly learn truth, safety, or usefulness. It learns to predict the preferences represented in its training data:

$r_{\theta}(x, y) \approx$ human preference for response y to prompt x .

Therefore the learned reward may reflect

- annotator preferences and disagreements,
- biases toward particular styles or levels of detail,
- incomplete coverage of possible prompts and behaviors.

The learned reward is a proxy for the behavior we actually want.

FROM REWARD MODEL TO REINFORCEMENT LEARNING

Now keep the trained reward model r_θ fixed and optimize the language model.

prompt $\rightarrow$ token $\rightarrow$ token $\rightarrow \dots \rightarrow$ response $\rightarrow$ reward

- π_ϕ^{RL} : language model being optimized.
- Given a prompt x , it samples a response $y \sim \pi_\phi^{\text{RL}}(\cdot \mid x)$.
- The fixed reward model assigns $r_\theta(x, y)$.
- PPO updates ϕ so that high-reward responses become more likely to be generated.

**Key difference from reward-model training:
the training responses are generated by the model being
optimized.**

THE RLHF OBJECTIVE

In InstructGPT, PPO optimizes the policy parameters ϕ :

$$\text{objective}(\phi) = \mathbb{E}_{(x,y) \sim D_{\pi_{\phi}^{\text{RL}}}} \left[r_{\theta}(x, y) - \beta \log \left(\frac{\pi_{\phi}^{\text{RL}}(y | x)}{\pi^{\text{SFT}}(y | x)} \right) \right].$$

- π_{ϕ}^{RL} : current policy; ϕ **is optimized**.
- π^{SFT} : fixed SFT reference model.
- r_{θ} : fixed trained reward model.
- $D_{\pi_{\phi}^{\text{RL}}}$: prompt–response samples generated under the current policy.

Unlike supervised learning, the sample distribution itself changes as ϕ changes.

► Ouyang et al., 2022

HOW TO READ THE RLHF OBJECTIVE

$$\mathbb{E}_{(x,y) \sim D_{\pi_{\phi}^{\text{RL}}}} \left[\underbrace{r_{\theta}(x, y)}_{\text{high reward}} - \underbrace{\beta \log \left(\frac{\pi_{\phi}^{\text{RL}}(y | x)}{\pi^{\text{SFT}}(y | x)} \right)}_{\text{stay close to SFT}} \right]$$

The expectation means: approximately average over a batch sampled from the **current RL policy**:

$$\mathbb{E}_{(x,y) \sim D_{\pi_{\phi}^{\text{RL}}}} [f(x, y)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i, y_i), \quad y_i \sim \pi_{\phi}^{\text{RL}}(\cdot | x_i).$$

- High reward encourages the policy to generate preferred responses.
- The log-ratio penalizes moving too far from the SFT model.

PPO-PTX: PRESERVE PRETRAINING CAPABILITIES

InstructGPT also mixes pretraining gradients into PPO:

$$\text{objective}(\phi) = \mathbb{E}_{(x,y) \sim D_{\pi_{\phi}^{\text{RL}}}} \left[r_{\theta}(x,y) - \beta \log \left(\frac{\pi_{\phi}^{\text{RL}}(y | x)}{\pi^{\text{SFT}}(y | x)} \right) \right] \\ + \gamma \mathbb{E}_{x \sim D_{\text{pretrain}}} [\log \pi_{\phi}^{\text{RL}}(x)] .$$


- First line: high reward while staying close to SFT.
- Second line: keep assigning high likelihood to pretraining text.
- For plain PPO, $\gamma = 0$.
- With $\gamma > 0$, the paper calls the model **PPO-ptx**.

► Ouyang et al., 2022

DOES PPO IMPROVE OVER SFT?

In the main InstructGPT human evaluation, PPO-based models were preferred over SFT.

- PPO win rate over SFT is roughly 0.6: clearly above chance-level 0.5.
- The improvement is **meaningful, but not dramatic**.
- PPO-ptx also reduces regression on pretraining benchmarks.



► Ouyang et al., 2022: InstructGPT

**Can one achieve similar gains
with a simpler
preference-optimization
method?**


Direct Preference Optimization

CAN WE OPTIMIZE PREFERENCES DIRECTLY?

Standard RLHF uses two separate optimization stages:

human preferences $\longrightarrow$ reward model $\longrightarrow$ RL / PPO.

Direct Preference Optimization (DPO): **Optimize the language model directly from the preference pairs** ► Rafailov et al., 2023: Direct Preference Optimization



- No separately trained reward model.
- No sampling from the policy inside an RL training loop.
- Train directly on a fixed preference dataset.

THE DPO OBJECTIVE

Training data are preference triples

$$(x, y_w, y_l),$$

where y_w is the winning and y_l the lower-ranked response.

DPO minimizes

$$\mathcal{L}_{\text{DPO}}(\theta) = -\mathbb{E}_{(x, y_w, y_l) \sim D} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \beta \log \frac{\pi_{\theta}(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right) \right].$$

- π_{θ} : policy being optimized.
- π_{ref} : fixed reference policy, typically the SFT model.
- β : inverse-temperature parameter; scales the effect of a given preference margin.
- D : fixed dataset of preference comparisons.

HOW TO READ THE DPO OBJECTIVE

Consider

$$\log \frac{\pi_{\theta}(y \mid x)}{\pi_{\text{ref}}(y \mid x)}.$$

It measures how much the new policy has changed the likelihood of response y relative to the reference model.

The DPO loss compares this change for the two responses:

$$\underbrace{\log \frac{\pi_{\theta}(y_w \mid x)}{\pi_{\text{ref}}(y_w \mid x)}}_{\text{preferred response}} - \underbrace{\log \frac{\pi_{\theta}(y_l \mid x)}{\pi_{\text{ref}}(y_l \mid x)}}_{\text{rejected response}}.$$

- Make the preferred response relatively more likely.
- Make the rejected response relatively less likely.
- Measure both changes relative to the fixed reference policy.

WHY CAN DPO ELIMINATE THE REWARD MODEL?

For KL-regularized (w.r.t. the reference policy) RLHF, ► Rafailov et al., 2023

$$r(x, y) = \beta \log \frac{\pi^*(y | x)}{\pi_{\text{ref}}(y | x)} + c(x).$$


DPO reverses this relation: **Each candidate policy defines the reward for which it would be optimal.**

Thus, for a preference pair,

$$r_{\theta}(x, y_w) - r_{\theta}(x, y_l) = \beta \left[\log \frac{\pi_{\theta}(y_w | x)}{\pi_{\text{ref}}(y_w | x)} - \log \frac{\pi_{\theta}(y_l | x)}{\pi_{\text{ref}}(y_l | x)} \right],$$

because the prompt-dependent constant cancels.

DOES DPO WORK?




- DPO matches or outperforms PPO in the experiments in the original paper.
- On TL;DR summarization: about 61% win rate for DPO vs. about 57% for PPO at their best reported settings.
- DPO achieves this without a separate reward model or an RL training loop.

Alignment Trade-offs

SAFETY TRAINING CAN OVER-REFUSE

Safety training should reduce harmful responses, but too much caution can make the model less useful.



Alignment requires a trade-off between helpfulness and harmlessness.

- Too permissive → harmful responses.
- Too cautious → harmless requests are refused.

REWARD HACKING: OPTIMIZING THE PROXY

The reward model is only a proxy for the behavior we actually want.

desired behavior $\longrightarrow$ human preferences $\longrightarrow$ learned reward

Example: hallucination

A response may be

fluent + confident + helpful-looking + **wrong**

yet still receive high reward if the reward model has learned to value the first three properties strongly.

high learned reward $\nrightarrow$ high true quality.

Optimizing a proxy can amplify its mistakes.

ALIGNMENT IS MULTI-OBJECTIVE

What counts as a “good” response is not a single objective.

helpful truthful harmless uncertain when appropriate

These goals can conflict, so post-training implicitly chooses **trade-offs** between them.

- Preferences depend on annotators and their instructions.
- Different users and cultures may prefer different behavior.
- There is therefore no completely value-neutral alignment target.

Alignment is not only an optimization problem: it also specifies whose preferences the model should follow.


Reasoning and Chain-of-Thought

CHAIN-OF-THOUGHT: THE ORIGINAL IDEA

Standard prompting often struggles on tasks that require multi-step reasoning.

Chain-of-thought prompting adds intermediate reasoning steps:

input $\longrightarrow$ reasoning steps $\longrightarrow$ answer.



- Originally introduced as a form of few-shot prompting.
- Demonstrations contain a reasoning trace and the final answer.

CHAIN-OF-THOUGHT: PROMPTING VS. TRAINING

The term “chain-of-thought” is now used more broadly.

- **Original CoT few-shot prompting:** reasoning examples are provided in the prompt.
- “Think step by step” without examples is not few-shot chain-of-thought prompting.
- Modern reasoning models can instead be **trained to generate reasoning traces**.

Reasoning traces need not come from examples in the prompt.

REASONING AS POST-TRAINING

Modern reasoning models use post-training to make reasoning behavior part of the policy itself.

- Supervised fine-tuning on reasoning examples.
- Reinforcement learning on generated reasoning trajectories.
- The model can produce a potentially long reasoning trace before giving the final answer.

Reasoning becomes learned behavior, not just a prompting trick.

► DeepSeek-AI, 2025

Query

The denominator of a fraction is 7 less than 3 times the numerator.
If the fraction is equivalent to $2/5$,
what is the numerator?

Correct reasoning trace

- | | |
|--------|------------------------------------|
| Step 1 | Let the numerator be x . |
| Step 2 | The denominator is $3x-7$. |
| Step 3 | We know that $x/(3x-7) = 2/5$. |
| Step 4 | Therefore, $5x = 6x-14$. |
| Step 5 | Finally, we get $x=14$. (Correct) |

REASONING ENABLES VERIFIABLE REWARDS

For some tasks, correctness can be checked automatically.

- **Mathematics:** compare the final answer with a known correct answer.
- **Code:** compile and run generated programs against test cases.

PROMPT

Let $b \geq 2$ be an integer. Call a positive integer n *b-eautiful* if it has exactly two digits when expressed in base b , and these two digits sum to $\sqrt{n}$. For example, 81 is 13-eautiful because $81 = \underline{6} \underline{3}_{13}$ and $6 + 3 = \sqrt{81}$. Find the least integer $b \geq 2$ for which there are more than ten *b-eautiful* integers.

Please reason step by step, and put your final answer within `\boxed{}`.

Evaluation

Parse the final answer within `\boxed{}` and use a rule-based grader to determine if it equals the ground truth. Round numerical values as needed, and use 'SymPy' to parse expressions.

The reward can come from an automatic verifier instead of human preferences.

► DeepSeek-AI, 2025

WHAT CAN REASONING TRAINING CHANGE?

During reasoning training, models can learn behaviors such as

- trying alternative strategies,
- self-verification and checking,
- reflection and correction,
- backtracking,
- producing longer reasoning traces for harder problems.

Question: If $a > 1$, then the sum of the real solutions of $\sqrt{a - \sqrt{a + x}} = x$ is equal to

Response: <think>

To solve the equation $\sqrt{a - \sqrt{a + x}} = x$, let's start by squaring both ...

$$(\sqrt{a - \sqrt{a + x}})^2 = x^2 \implies a - \sqrt{a + x} = x^2.$$

Rearrange to isolate the inner square root term:

$$(a - x^2)^2 = a + x \implies a^2 - 2ax^2 + (x^2)^2 = a + x \implies x^4 - 2ax^2 - x + (a^2 - a) = 0$$

...

Wait, wait. Wait. That's an aha moment I can flag here.

Let's reevaluate this step-by-step to identify if the correct sum can be ...


We started with the equation:

$$\sqrt{a - \sqrt{a + x}} = x$$

First, let's square both sides:

$$a - \sqrt{a + x} = x^2 \implies \sqrt{a + x} = a - x^2$$

Next, I could square both sides again, treating the equation: ...



A PLAUSIBLE REASONING TRACE

The model answers correctly and then explains its answer with a familiar step-by-step algorithm.

Human: Answer in one word. What is $36+59$?


Assistant: 95

Human: Briefly, how did you get that?

Assistant: I added the ones ($6+9=15$), carried the 1, then added the tens ($3+5+1=9$), resulting in 95.

But did the model actually compute the answer this way?

REASONING TRACES ARE NOT NECESSARILY FAITHFUL



Interpretability analysis suggests that the internal computation can differ substantially from the verbalized algorithm. ▶ Anthropic, 2025

POST-TRAINING: MAIN TAKEAWAYS

Pretrained LM $\longrightarrow$ SFT $\longrightarrow$ Preference Optimization

- **SFT:** train on desired instruction–response pairs.
- **Reward model:** train on human preference comparisons between responses.
- **PPO / RLHF:** optimize a policy that generates high-reward responses while staying close to the SFT reference model.
- **DPO:** optimize the policy directly from preference pairs, without fitting a separate reward model or running an RL loop.
- **Alignment has trade-offs:** imperfect objectives can cause reward hacking and over-refusal.
- **Reasoning can be post-trained:** verifiable feedback can encourage checking, reflection, and backtracking.